

Multi-Agent Financial Analysis System - Design Document

Version: 1.0 | **Date:** 5 October 2025 | **Type:** LLD

1. System Overview

A multi-agent AI system that analyzes US public companies by combining SEC 10-K filings with real-time market data. Uses specialized agents coordinated by an orchestrator, advanced RAG for document retrieval, and dual-memory architecture for personalization.

Core Stack: phidata (orchestration) | OpenAI GPT-4o | ChromaDB (vectors) | yfinance (market data) | SQLite (persistence)

2. Agent Roles

Orchestrator (orchestrator.py)

Purpose: Multi-agent coordinator

Key Function: `create_plan()` decomposes queries into DAG of agent tasks, `execute_plan()` runs agents sequentially based on dependencies

Example Plan:

```
{"steps": [  
  {"agent": "sec_researcher", "dependencies": []},  
  {"agent": "market_data", "dependencies": []},  
  {"agent": "analyst", "dependencies": ["sec_researcher", "market_data"]},  
  {"agent": "auditor", "dependencies": ["analyst"]}  
]}
```

SEC Researcher (sec_researcher.py)

Purpose: 10-K filing analyst using RAG

Key Functions: `analyze_risks(ticker, years)` → risk factors by year | `get_financial_trends()` → revenue/margin/cash flow

RAG Strategy: Query decomposition → Hybrid search (dense + sparse) → Re-ranking → Top 5 results

Market Data Agent (market_data.py)

Purpose: Real-time metrics via Yahoo Finance

Key Function: `get_comprehensive_data(ticker)` → {price, market_cap, P/E, beta, 52-week range, timestamp}

Validation: Verifies $\text{market_cap} \approx \text{price} \times \text{shares}$ ($\pm 5\%$ tolerance)

Financial Analyst (analyst.py)

Purpose: Report synthesizer

Key Function: `synthesize(sec_data, market_data, ticker)` → Markdown report with sections: Executive Summary, Market Metrics, Business Overview, Risk Factors, Financial Trends, Investment View

Citation: Tracks all sources via CitationTracker

Auditor (auditor.py)

Purpose: Quality assurance

Key Function: `verify(report, citations, market_data)` → {citation_check, numeric_check, claim_check, overall_confidence}

Validation Rules: All factual claims need citations | Market cap calculations within 5% | Confidence = (passed checks) / 3

3. RAG Pipeline Architecture

File: `rag_pipeline.py` | **Class:** AdvancedRAGPipeline

4-Stage Retrieval Process

Stage 1: Query Decomposition

- Uses GPT-3.5-turbo to break complex queries into 2-3 sub-queries
- Fallback: Original query if LLM fails or returns invalid JSON

Stage 2: Hybrid Search (Top 20)

- **Dense:** Sentence-BERT embeddings (768-dim) + cosine similarity
- **Sparse:** BM25 keyword matching (exact terms/numbers)
- **Fusion:** RRF (Reciprocal Rank Fusion) score = $1/(60+\text{rank})$ merges both

Stage 3: Re-ranking (Top 5)

- Cross-encoder model: cross-encoder/ms-marco-MiniLM-L-6-v2
- Scores each passage for relevance
- More accurate than bi-encoders but slower (only used on top-20)

Stage 4: Cross-Document Analysis

- `compare_across_years(query, [2020-2024])` → Results grouped by year
- Enables trend analysis across multiple 10-K filings

4. Memory Schema

Global Memory (Persistent) - `memory.py`

Storage: SQLite (`global_memory.db`) + ChromaDB

Tables:

```
user_preferences (  
    user_id, risk_taxonomy JSON, writing_style,  
    preferred_kpis JSON, version, created_at, ttl_expires  
)
```

```
analysis_history (  
    ticker, analysis_date, summary, key_metrics JSON, embedding BLOB  
)
```

Use Cases: Personalized reports | Similar company search | Analysis history tracking

TTL: 365 days for preferences.